

Introduction to Data Management

SQL Basics and Joins

Paul G. Allen School of Computer Science and Engineering
University of Washington, Seattle

- **Relational data model**
 - Schema+instance+query language
- **Query language: SQL**
 - Create tables
 - Retrieve records from tables
 - Declare keys and foreign keys

Keys

Key

A **Key** is one or more attributes that, in aggregate, uniquely identify a row.

Key

<u>cname</u>	country	no_employees	for_profit
GizmoWorks	USA	20000	True
Canon	Japan	50000	True
Hitachi	Japan	30000	True
HappyCam	Canada	500	False

Keys

Key

A **Key** is one or more attributes that, in aggregate, uniquely identify a row.

Key

Not a key

<u>cname</u>	country	no_employees	for_profit
GizmoWorks	USA	20000	True
Canon	Japan	50000	True
Hitachi	Japan	30000	True
HappyCam	Canada	500	False

Keys

Key

A **Key** is one or more attributes that, in aggregate, uniquely identify a row.

Key

Not a key

Is this a key?

<u>cname</u>	country	no_employees	for_profit
GizmoWorks	USA	20000	True
Canon	Japan	50000	True
Hitachi	Japan	30000	True
HappyCam	Canada	500	False

Keys

Key

A **Key** is one or more attributes that, in aggregate, uniquely identify a row.

Key

Not a key

Is this a key?

No: future updates to the database may create duplicate no_employees


<u>cname</u>	country	no_employees	for_profit
GizmoWorks	USA	20000	True
Canon	Japan	50000	True
Hitachi	Japan	30000	True
HappyCam	Canada	500	False

Multi-attribute Key

Key

A **Key** is one or more attributes that, in aggregate, uniquely identify a row.

Key = fName, lName
(what does this mean?)

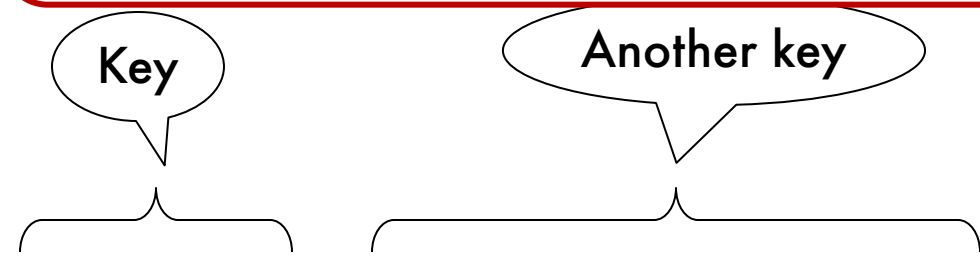


<u>fName</u>	<u>lName</u>	Income	Department
Alice	Smith	20000	Testing
Alice	Thompson	50000	Testing
Bob	Thompson	30000	SW
Carol	Smith	50000	Testing

Multiple Keys

Key

A **Key** is one or more attributes that, in aggregate, uniquely identify a row.



<u>SSN</u>	fName	IName	Income	Department
111-22-3333	Alice	Smith	20000	Testing
222-33-4444	Alice	Thompson	50000	Testing
333-44-5555	Bob	Thompson	30000	SW
444-55-6666	Carol	Smith	50000	Testing

We can choose one key and designate it as primary key
E.g.: primary key = SSN

Foreign Key

Company

<u>cname</u>	country	no_employees	for_profit
Canon	Japan	50000	Y
Hitachi	Japan	30000	Y

Country

<u>cname</u>	population
USA	320M
Japan	127M

Foreign Key

Foreign key to
Country.cname

Company

<u>cname</u>	country	no_employees	for_profit
Canon	Japan	50000	Y
Hitachi	Japan	30000	Y

Country

<u>cname</u>	population
USA	320M
Japan	127M

Foreign Key

Foreign key to
Country.cname

Company

<u>cname</u>	country	no_employees	for_profit
Canon	Japan	50000	Y
Hitachi	Japan	30000	Y

Country

<u>cname</u>	population
USA	320M
Japan	127M

Foreign Key

A **Foreign Key** is one or more attributes that, in aggregate, uniquely identify a row in *another table*.

Review

- Tables are NOT ordered
 - they are sets or multisets (bags)

Review

- Tables are NOT ordered
 - they are sets or multisets (bags)
- Tables are FLAT
 - No nested attributes
 - We say that the table is in **First Normal Form**

Review

- Tables are NOT ordered
 - they are sets or multisets (bags)
- Tables are FLAT
 - No nested attributes
 - We say that the table is in **First Normal Form**
- Tables DO NOT prescribe how they are implemented / stored on disk

Review

- Tables are NOT ordered
 - they are sets or multisets (bags)
- Tables are FLAT
 - No nested attributes
 - We say that the table is in **First Normal Form**
- Tables DO NOT prescribe how they are implemented / stored on disk
- **Physical data independence**

SQL

- **Structured Query Language**
- Most widely used language to query relational data
- One of the many languages for querying relational data
- A **declarative** programming language

Note

SQL is a huge language. We cover in class the important concepts. You may have to google for some simple command that we don't cover in class.

Simplest variant: three clauses:

- **SELECT:** which attributes we want
- **FROM:** which table(s) we query
- **WHERE:** predicate saying which rows we want

URLs

ID	URL	Title	NumVisits
5	cs.washington.edu	Welcome to Paul G.	13
6	google.com	Google	2
34	washington.zoom.us	Video Conf...	4
37	canvas.uw.edu	Dashboard	6

URLs

ID	URL	Title	NumVisits
5	cs.washington.edu	Welcome to Paul G.	13
6	google.com	Google	2
34	washington.zoom.us	Video Conf...	4
37	canvas.uw.edu	Dashboard	6

```
SELECT URL, Title
FROM URLs
WHERE NumVisits > 5;
```

URLs

ID	URL	Title	NumVisits
5	cs.washington.edu	Welcome to Paul G.	13
6	google.com	Google	2
34	washington.zoom.us	Video Conf...	4
37	canvas.uw.edu	Dashboard	6

```
SELECT URL, Title
FROM URLs
WHERE NumVisits > 5;
```

Answer:

URL	Title
cs.washington.edu	Welcome to Paul G.
canvas.uw.edu	Dashboard

Demo

Demo will show the following:

- CREATE TABLE
- INSERT, DELETE
- NULL
- More SELECT:
 - SELECT *
 - SELECT DISTINCT
 - Discuss the notions of filter(selection), projection
 - LIKE
 - ORDER BY

Joins in SQL

Product

<u>pname</u>	price	category	manufacturer
MultiTouch	199.99	gadget	Canon
SingleTouch	49.99	photography	Canon
Gizom	50	gadget	GizmoWorks
SuperGizmo	250.00	gadget	GizmoWorks

Company

<u>cname</u>	country
GizmoWorks	USA
Canon	Japan
Hitachi	Japan

cname is a key in Company
pname is a key in Product
manufacturer is a foreign to Company.cname

Joins in SQL

Product

<u>pname</u>	price	category	manufacturer
MultiTouch	199.99	gadget	Canon
SingleTouch	49.99	photography	Canon
Gizom	50	gadget	GizmoWorks
SuperGizmo	250.00	gadget	GizmoWorks

Company

<u>cname</u>	country
GizmoWorks	USA
Canon	Japan
Hitachi	Japan

Retrieve all Japanese products that cost < \$150

Joins in SQL

Product

<u>pname</u>	price	category	manufacturer
MultiTouch	199.99	gadget	Canon
SingleTouch	49.99	photography	Canon
Gizom	50	gadget	GizmoWorks
SuperGizmo	250.00	gadget	GizmoWorks

Company

<u>cname</u>	country
GizmoWorks	USA
Canon	Japan
Hitachi	Japan

Retrieve all Japanese products that cost < \$150

```
SELECT pname
FROM Product, Company
```

Joins in SQL

Product

<u>pname</u>	price	category	manufacturer
MultiTouch	199.99	gadget	Canon
SingleTouch	49.99	photography	Canon
Gizom	50	gadget	GizmoWorks
SuperGizmo	250.00	gadget	GizmoWorks

Company

<u>cname</u>	country
GizmoWorks	USA
Canon	Japan
Hitachi	Japan

Retrieve all Japanese products that cost < \$150

```
SELECT pname
FROM Product, Company
WHERE country = 'Japan'
AND price < 150
```

Joins in SQL

Product

<u>pname</u>	price	category	manufacturer
MultiTouch	199.99	gadget	Canon
SingleTouch	49.99	photography	Canon
Gizom	50	gadget	GizmoWorks
SuperGizmo	250.00	gadget	GizmoWorks

Company

<u>cname</u>	country
GizmoWorks	USA
Canon	Japan
Hitachi	Japan

Retrieve all Japanese products that cost < \$150

```
SELECT pname
FROM Product, Company
WHERE country = 'Japan'
AND price < 150
AND manufacturer = cname
```

Joins in SQL

Product

<u>pname</u>	price	category	manufacturer
MultiTouch	199.99	gadget	Canon
SingleTouch	49.99	photography	Canon
Gizom	50	gadget	GizmoWorks
SuperGizmo	250.00	gadget	GizmoWorks

Company

<u>cname</u>	country
GizmoWorks	USA
Canon	Japan
Hitachi	Japan

Retrieve all Japanese products that cost < \$150

```
SELECT pname
FROM Product, Company
WHERE country = 'Japan'
AND price < 150
```

What does this return?

Joins in SQL

Product

<u>pname</u>	price	category	manufacturer
MultiTouch	199.99	gadget	Canon
SingleTouch	49.99	photography	Canon
Gizom	50	gadget	GizmoWorks
SuperGizmo	250.00	gadget	GizmoWorks

Company

<u>cname</u>	country
GizmoWorks	USA
Canon	Japan
Hitachi	Japan

Retrieve all Japanese products that cost < \$150

```
SELECT pname
FROM Product, Company
WHERE country = 'Japan'
AND price < 150
```

What does this return?

ALL products < \$150

Demo

Demo will show

- foreign keys
- Joins
- cartesian product

Discussion

To join two tables:

- Write those table in the FROM clause
- Write explicitly the join condition
- Records that do not join with anything in the other table are not included