

SQL for Business Intelligence

Turning Data into Insights

0. Table of Contents

1. Introduction
2. What is a Business Intelligence?
3. Introduction to SQL in BI
4. Setting up Your BI Environment
5. Basic SQL Queries for BI

-- SELECT Statements

-- Filtering Data

-- Aggregating Data

6. Advanced SQL Techniques for BI

-- JOIN Operations Subqueries

-- Window Functions

7. Creating Interactive Reports
8. Data Visualization with SQL
9. Security and Access Control
10. Performance Optimization
11. Real-World BI Examples
12. Conclusion

1. Introduction

1.1 Question:

What kind of world we live in?

1.2 Answer:

1. In the data-driven world we live in,
2. businesses are constantly collecting vast amounts of data.

1.3 BUT

But what's the use of all that data if you can't turn it into actionable insights?

This is where Business Intelligence (BI) comes into play.

At the core of BI is SQL (Structured Query Language), a powerful tool for querying and analyzing data. In this article, we'll explore how SQL is used for Business Intelligence with practical code examples.

2. What is a Business Intelligence?

1. What is business intelligence?

Business intelligence combines

- business analytics,
- data mining,
- data visualization,
- data tools and infrastructure, and
- best practices

to help organizations make more data-driven decisions.

2. What Is Business Intelligence (BI)?

Business intelligence refers to the technology that enables businesses to

1. organize, analyze, and
 2. contextualize business data
- from around the company.

BI includes

1. multiple tools and techniques
2. to transform raw data into
3. meaningful and actionable information.

2.1 Examples of Business Intelligence

1. Why are sales dropping in a particular region?
2. What factors contribute to high profits in a specific year?
3. Which employee is performing better in this quarter?
4. What are the top-10 products sold in California
5. How many orders were NOT FILLED IN last week?
6. What was the totals sales of Nike shoes in 58 counties of California for the past 3 years?
7. Which are the top-5 states for selling Electric bikes?

2.2 Business Intelligence Main Parts

BI systems have 4 main parts:

1. A data warehouse stores company information from a variety of sources in a centralized and accessible location.
2. Business analytics or data management tools mine and analyze data in the data warehouse.
3. Business performance management (BPM) tools monitor and analyze progress towards business goals.
4. A user interface (usually an interactive dashboard with data visualization reporting tools) provides quick access the information.

2.3 Companies using Business INtelligence

- Starbucks

BI lets them provide loyalty card programs and a mobile app to track individual purchase data. It also helps predict purchases and send personalized offers via app or email to entice customers and boost sales.

Starbucks carefully analyses a number of data points to identify what might make a successful store location. This includes:

1. Consumer demographics
2. Population density
3. Average income levels
4. Traffic patterns
5. Public transport hubs
6. Types of businesses in the location under consideration.

- Netflix

As the largest subscription streaming service, Netflix leverages business intelligence and data science to gain new subscribers and enhance user experiences.

BI solutions create personalized watch lists and compelling artwork while persuading users to watch their favorite genres. The recommendation system promotes content to the right audience at the right time.

- Uber

This California-based transportation company uses business intelligence to monitor traffic conditions, journey times and durations, driver availability and customer demand in real time.

It leverages surge pricing and a rating system as a part of its business model to manage fares and maintain trust between drivers and customers.

- Walmart

Walmart uses business intelligence to gauge customer behavior online and in-store based on customer purchases. It helps them make data-driven decisions around assortments, inventory and merchandising, resources, promotions and more to drive profits. Data-based apps shine a light on customer preferences and curate specific loyalty programs for different segments.

- Amazon

Amazon is a retail giant with more than 100 million prime customers worldwide and over 12 million products for sale.

The global marketplace is all about concentrating huge amounts of data that should be processed quickly and effectively to ensure all personalized user recommendations, targeted marketing, and pricing optimizations are up to date. All these analytic efforts result in the extended collection of various metrics about every customer and help bring the user experience to such a new, high-quality level that makes Amazon a leading e-retailer in the US.

Amazon uses BI systems to optimize its logistics and streamline the workflow of the supply chain. Utilizing modern BI tools, Amazon speeds up order deliveries thanks to predictive analysis, choosing the best shipping schedules, routes, and groupings of products.

2.4 What are the Techniques for Business Intelligence

1. Analytics:

Extracts and identifies trends, patterns and correlations from historical data to drive decisions.

2. Dashboards:

Collects and displays intuitive visualizations, KPIs and metrics to monitor business performance.

3. Data Mining:

Employs statistical and machine learning techniques to unearth anomalies, patterns and correlations within large datasets to predict outcomes.

4. ETL:

Extracts data from multiple sources, cleans it to improve quality and loads it into systems for access and manipulation.
-> Build a Data Warehouse

5. OLAP:

Performs complex analytical queries on Data Warehouses to solve problems from different perspectives.

6. Predictive Modeling:

Uses machine learning and data mining techniques to predict future outcomes by analyzing historical data.

3. Introduction to SQL in BI

> Business Intelligence (BI) involves the use of data analysis tools and techniques to transform raw data into valuable insights.

> SQL, as a querying language, plays a pivotal role in BI.

> SQL allows BI professionals to interact with

databases, retrieve relevant data, and perform various data transformations.

4. Setting up Your BI Environment

Before we dive into SQL for BI, you need the right tools and environment. Typically, this involves setting up a data warehouse or connecting to existing databases.

Popular BI tools like Tableau, Power BI, and Looker often integrate seamlessly with SQL databases.

4.1 Code Example: Connecting to a Database

```
-- Connect to a SQL Server database
USE employee_db;

-- Select data from a table
SELECT FirstName, LastName
FROM Employees
WHERE Department = 'Sales';
```

4.2 Code Example: Find Number of Employees per Department

```
-- Connect to a SQL Server database
USE employee_db;

-- Select data from a table and Group them
SELECT Department, count(*) as total
FROM Employees
GROUP BY Department;
```

5. Basic SQL Queries for BI

5.1 SELECT Statements

The SELECT statement is the workhorse of SQL. It allows you to retrieve data from one or more tables.

In BI, this is the starting point for fetching data.

5.2 Filtering Data

Filtering data is crucial for BI. You want to focus on the data that's relevant to your analysis.

SQL's WHERE clause helps you do that.

5.3 Aggregating Data

Aggregation functions like SUM , AVG , and COUNT are essential for summarizing data.

These functions are used extensively in creating BI reports.

5.4 Code Example: Basic SQL Queries

- Retrieve all product names and their prices

```
SELECT ProductName, Price  
FROM Products;
```

- Filter orders for a specific customer


```
SELECT OrderID, OrderDate
FROM Orders
WHERE CustomerID = 'C2007';
```

- Calculate total sales for each year

```
SELECT YEAR(OrderDate) AS Year,
       SUM(TotalAmount) AS TotalSales
FROM Orders
GROUP BY YEAR(OrderDate);
```

6. Advanced SQL Techniques for BI

6.1 JOIN Operations

In BI, data often resides in multiple tables. SQL's JOIN operations (e.g., INNER JOIN, LEFT JOIN, ...) allow you to combine data from different tables.

```
-- books: (book_id, book_title, author_id)
-- authors:(author_id, first_name, last_name)
--
SELECT b.book_id, b.book_title, a.first_name, a.last_name
FROM books b
     INNER JOIN authors a ON b.author_id = a.author_id
ORDER BY b.book_id;
```

6.2 Subqueries

Subqueries (or nested queries) are powerful tools for BI. They let you create complex queries step by step.

```
SELECT * FROM CUSTOMERS
WHERE ID IN (
    SELECT ID
    FROM CUSTOMER_DETAILS
    WHERE SALARY > 4500
```

```
);
```

6.3 GROUP BY year WITH ROLLUP;

```
mysql> SELECT * from product_sales;
```

year	country	product	profit
2000	USA	PC	700
2000	USA	IPAD	800
2000	USA	MAC	900
2000	CANADA	PC	900
2000	CANADA	IPAD	300
2001	USA	PC	200
2001	USA	IPAD	300
2001	USA	MAC	400
2001	CANADA	PC	600
2001	CANADA	IPAD	400

10 rows in set (0.00 sec)

```
mysql> SELECT year, country, SUM(profit) AS profit
-> FROM product_sales
-> GROUP BY year, country;
```

year	country	profit
2000	USA	2400
2000	CANADA	1200
2001	USA	900
2001	CANADA	1000

4 rows in set (0.00 sec)

```
mysql> SELECT year, country, SUM(profit) AS profit
-> FROM product_sales
-> GROUP BY year, country with ROLLUP
-> order by year;
```

year	country	profit
NULL	NULL	5500
2000	CANADA	1200
2000	USA	2400
2000	NULL	3600
2001	CANADA	1000

<-- 5500 = 2400+1200+900+1000
<-- 3600 = 2400+1200

2001	USA	900	
2001	NULL	1900	<-- 1900 = 1000+900

7 rows in set (0.01 sec)

-- NOTE: The NULL value in the (year, country) columns
 -- identifies the grand total super-aggregate line.

6.4 GROUP by CUBE

GROUP BY CUBE is an extension of the GROUP BY clause similar to GROUP BY ROLLUP. In addition to producing all the rows of a GROUP BY ROLLUP, GROUP BY CUBE adds all the cross-tabulations rows. Sub-total rows are rows that further aggregate whose values are derived by computing the same aggregate functions that were used to produce the grouped rows.

[6315 group by cube rollup grouping sets sql server.006 2] |

https://www.mssqltips.com/tipimages2/6315_group-by-cube-rollup-grouping-sets-sql-server.006-2.png

A CUBE grouping is equivalent to a series of grouping sets and is essentially a shorter specification. The N elements of a CUBE specification correspond to 2^N GROUPING SETS.

[1*Gq5Ec6OsqrI0UTgdxAhemQ] |

https://miro.medium.com/v2/resize:fit:1093/1*Gq5Ec6OsqrI0UTgdxAhemQ.png

Snowflake Examples

```
-- Create some tables and insert some rows.
CREATE TABLE products (product_ID INTEGER, wholesale_price REAL);

INSERT INTO products (product_ID, wholesale_price) VALUES
  (1, 1.00),
  (2, 2.00);

CREATE TABLE sales (product_ID INTEGER, retail_price REAL,
  quantity INTEGER, city VARCHAR, state VARCHAR);
INSERT INTO sales (product_id, retail_price, quantity, city, state) VALUES
  (1, 2.00, 1, 'SF', 'CA'),
  (1, 2.00, 2, 'SJ', 'CA'),
  (2, 5.00, 4, 'SF', 'CA'),
  (2, 5.00, 8, 'SJ', 'CA');
```

```
(2, 5.00, 16, 'Miami', 'FL'),
(2, 5.00, 32, 'Orlando', 'FL'),
(2, 5.00, 64, 'SJ', 'PR');
```

Run a cube query that shows profit by city, state, and total across all states. The example below shows a query that has three `levels`:

- Each city.
- Each state.
- All revenue combined.

This example uses `ORDER BY state, city NULLS LAST` to ensure that each state's rollup comes immediately after all of the cities in that state, and that the final rollup appears at the end of the output.

```
SELECT state, city,
       SUM((s.retail_price - p.wholesale_price) * s.quantity) AS profit
FROM products AS p, sales AS s
WHERE s.product_ID = p.product_ID
GROUP BY CUBE (state, city)
ORDER BY state, city NULLS LAST;
```

STATE	CITY	PROFIT
CA	SF	13
CA	SJ	26
CA	NULL	39
FL	Miami	48
FL	Orlando	96
FL	NULL	144
PR	SJ	192
PR	NULL	192
NULL	Miami	48
NULL	Orlando	96
NULL	SF	13
NULL	SJ	218
NULL	NULL	375

6.4 Window Functions

Window functions provide a way to perform calculations across a set of rows related to the current row.

This is extremely valuable for BI analytics.



Unlike aggregate functions, window functions do not collapse rows.

[aggregate vs window functions] | <https://learnsql.com/blog/sql-window-functions-cheat-sheet/aggregate-vs-window-functions.png>

6.5 Window Function Concepts and Syntax

This section describes how to use window functions. Examples use the same sales information data set.

```
mysql> SELECT *  
        FROM sales  
        ORDER BY country, year, product;
```

year	country	product	profit
2000	Finland	Computer	1500
2000	Finland	Phone	100
2001	Finland	Phone	10
2000	India	Calculator	75
2000	India	Calculator	75
2000	India	Computer	1200
2000	USA	Calculator	75
2000	USA	Computer	1500
2001	USA	Calculator	50
2001	USA	Computer	1500
2001	USA	Computer	1200
2001	USA	TV	150
2001	USA	TV	100

A window function performs an aggregate-like operation on a set of query rows. However, whereas an aggregate operation groups query rows into a single result row, a window function produces

a result for each query row:

The row for which function evaluation occurs is called the current row.

The query rows related to the current row over which function evaluation occurs comprise the window for the current row.

For example, using the sales information table, these two queries perform aggregate operations that produce a single global sum for all rows taken as a group, and sums grouped per country:

```
mysql> SELECT SUM(profit) AS total_profit
        FROM sales;
```

total_profit
7535

```
mysql> SELECT country, SUM(profit) AS country_profit
        FROM sales
        GROUP BY country
        ORDER BY country;
```

country	country_profit
Finland	1610
India	1350
USA	4575

By contrast, window operations do NOT collapse groups of query rows to a single output row. Instead, they produce a result for each row. Like the preceding queries, the following query uses SUM(), but this time as a window function:

```
mysql> SELECT
        year, country, product, profit,
```

```
SUM(profit) OVER() AS total_profit,
SUM(profit) OVER(PARTITION BY country) AS country_profit
FROM sales
ORDER BY country, year, product, profit;
```

year	country	product	profit	total_profit	country_profit
2000	Finland	Computer	1500	7535	1610
2000	Finland	Phone	100	7535	1610
2001	Finland	Phone	10	7535	1610
2000	India	Calculator	75	7535	1350
2000	India	Calculator	75	7535	1350
2000	India	Computer	1200	7535	1350
2000	USA	Calculator	75	7535	4575
2000	USA	Computer	1500	7535	4575
2001	USA	Calculator	50	7535	4575
2001	USA	Computer	1200	7535	4575
2001	USA	Computer	1500	7535	4575
2001	USA	TV	100	7535	4575
2001	USA	TV	150	7535	4575

Each window operation in the query is signified by inclusion of an OVER clause that specifies how to partition query rows into groups for processing by the window function:

The first OVER clause is empty, which treats the entire set of query rows as a single partition. The window function thus produces a global sum, but does so for each row.

The second OVER clause partitions rows by country, producing a sum per partition (per country). The function produces this sum for each partition row.

6.4 Code Example: Advanced SQL Queries

- Combine customer and order information:

```
SELECT Customers.CustomerName, Orders.OrderID, Orders.OrderDate
FROM Customers
INNER JOIN Orders ON Customers.CustomerID = Orders.CustomerID;
```

- Find customers with high total purchases:

```
SELECT CustomerName
FROM Customers
WHERE CustomerID IN
(
    SELECT CustomerID
    FROM Orders
    GROUP BY CustomerID
    HAVING SUM(TotalAmount) > 10000
);
```

- Calculate a moving average of sales

```
SELECT OrderDate,
    TotalAmount,
    AVG(TotalAmount) OVER (ORDER BY OrderDate) AS MovingAvg
FROM Orders;
```

7. Creating Interactive Reports

BI reports are interactive, allowing users to explore data. SQL can be used to create parameterized queries for dynamic reports.

7.1 Code Example: Creating Parameterized Query

```
-- Create a parameterized query for sales by region
DECLARE @Region VARCHAR(255) = 'West';

SELECT Region, ProductName, SUM(Quantity) AS TotalQuantity
FROM Sales
WHERE Region = @Region
GROUP BY Region, ProductName;
```

8. Data Visualization with SQL

BI isn't just about numbers;

it's about SQL and Visuals.

SQL can be used to retrieve data for visualization tools like charts and dashboards.

8.1 Code Example: Retrieving Data for Visualization

```
-- Retrieve data for a bar chart of monthly sales
SELECT DATE_FORMAT(OrderDate, '%Y-%m') AS Month,
       SUM(TotalAmount) AS MonthlySales
FROM Orders
GROUP BY Month
ORDER BY Month;
```

9. Security and Access Control

In a BI environment, data security is paramount. SQL can enforce access controls to ensure only authorized users can access certain data.

10. Performance Optimization

BI often deals with large datasets. SQL optimization techniques, like indexing and query tuning, can significantly boost performance.

11. Real-World BI Examples

Let's look at some real-world scenarios where SQL is used in BI:

1. Sales Analysis: SQL helps analyze sales data to identify trends, topselling products, and regions with high demand.

2. Customer Segmentation: Segmenting customers based on buying behavior using SQL's aggregation functions.
3. Financial Reporting: Generating financial reports like balance sheets and income statements using SQL.
4. Inventory Management: Using SQL to optimize inventory levels and avoid overstocking or understocking.
5. Marketing Campaign Analysis: Measuring the effectiveness of marketing campaigns through SQL queries.

12. Conclusion

1. SQL is the backbone of Business Intelligence.
2. SQL empowers professionals to access, manipulate, and analyze data, turning it into valuable insights.
3. Whether you're a BI analyst, data engineer, or business leader, understanding SQL is essential for making data-driven decisions.
4. So, roll up your sleeves and start querying — your data has stories to tell.
5. This article has provided an overview of how SQL is used in Business Intelligence.
6. To become proficient, practice is key. So, set up your BI environment, load some data, and start exploring. The world of data-driven insights awaits!

References

1. [SQL for Business Intelligence: Turning Data into Insights](#)
2. [SQL Window Functions Cheat Sheet](#)
3. [GROUP BY Modifiers](#)
4. [SQL Cube](#)
5. [Group By in SQL Server with CUBE, ROLLUP and GROUPING SETS Examples](#)